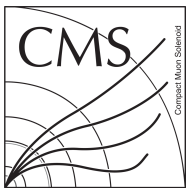
Internal review draft – 22 October 2026
No CMS DP number assigned



The Compact Muon Solenoid Experiment
CMS Performance Note



Dual-threshold time-walk calibration for CMS BTL prototypes

CMS

cms-dpg-conveners-mtd@cern.ch

Abstract

A charge-dependent time-walk model is studied for CMS barrel timing layer (BTL) prototype channels using two discriminator thresholds. Their current separation divided by the measured crossing-time separation provides a finite-interval pulse-slope proxy. A linear profile fit describes this proxy as a function of the measured charge, and an inverse-slope function models the timing correction with one amplitude parameter per channel.

Slope measurements from run 4076 are compared with independent microchannel-plate (MCP) reference calibrations from run 4399. The two runs contain different events. A free constant offset and the correction amplitude are fitted only where both measurements have populated charge support. The resulting agreement is a calibration-shape comparison, not an independent measurement of timing resolution.

The draft also examines channel-to-channel variations and an exploratory prediction of the correction amplitude from internal pulse information. Such a prediction could enable operation without an event-by-event external time reference, but its transferability and full uncertainty remain to be established.

Draft v1 for scientific discussion. Not approved for public release.
The following slides retain explicitly marked review items.

Dual-threshold time-walk calibration for CMS BTL prototypes

CMS

cms-dpg-conveners-mtd@cern.ch

Discussion draft; not approved for public release.

Abstract

A charge-dependent time-walk model is studied for CMS barrel timing layer (BTL) prototype channels using two discriminator thresholds. Their current separation divided by the measured crossing-time separation provides a finite-interval pulse-slope proxy. A linear profile fit describes this proxy as a function of the measured charge, and an inverse-slope function models the timing correction with one amplitude parameter per channel.

Slope measurements from run 4076 are compared with independent microchannel-plate (MCP) reference calibrations from run 4399. The two runs contain different events. A free constant offset and the correction amplitude are fitted only where both measurements have populated charge support. The resulting agreement is a calibration-shape comparison, not an independent measurement of timing resolution.

The draft also examines channel-to-channel variations and an exploratory prediction of the correction amplitude from internal pulse information. Such a prediction could enable operation without an event-by-event external time reference, but its transferability and full uncertainty remain to be established.

Two complementary runs, not the same events

Sample	Available information	Role in this study
Run 4076	Valid dual-threshold timing and charge	Measure the pulse-slope proxy $S(Q)$
Run 4399	MCP-associated reference timing and charge	Measure the empirical walk $W_{\mathrm{MCP}}(Q)$

The runs contain different particles and are never matched event by event. The comparison transfers a channel-response model between runs.

The upstream and downstream modules each contain bars read out at two ends. Detector/readout context is given in Refs. [1, 2]. The channel/bar mapping determines the coincidence pairs; the same channel number is compared across the two runs.

Review N2-1: Document gain, bias, thresholds, temperature, firmware and charge calibration to justify cross-run transfer. The run-4399 reference is the clock-corrected version now also documented in Note1.

From two crossing times to a walk model

For thresholds $I_1 < I_2$ and their valid crossing times t_1, t_2, define

$$S = \frac{I_2 - I_1}{t_2 - t_1} \quad [\mu\mathrm{A/ns}], \qquad \overline{S}(Q) \simeq c_0 + c_1 Q.$$

S is a finite-interval rising-edge slope proxy. It is not the current difference ΔI itself and is not necessarily the instantaneous derivative at the first threshold.

$$W_Q(Q) = \frac{A}{c_0 + c_1 Q}, \qquad t_{\mathrm{corr}} = t_{\mathrm{measured}} - W_Q(Q).$$

The amplitude A (AWC) has units of current and is constant for a channel and operating configuration. The correction varies event by event through Q.

Review N2-2: Freeze the safe-t_2 decoding, threshold conversion and physical timing cuts. This is an effective pulse model, not an exact waveform identity.

Selection and profile construction

- Run 4076: require exactly the two hits of the target bar in its module. Hits in the other module are not constrained; channel 192 is ignored.
- Retain valid rising-edge timing and the selected prompt-slope population, removing unphysical and slow-slope events before profiling.
- Retain charge above the lower half-maximum boundary of the fitted Landau–Gaussian peak. The selected data extend to the right of the peak.
- Form ten charge bins. Plot the mean S and the hit standard deviation (SD); use the standard error of the mean (SEM) for the linear-fit weights.
- Bootstrap the selected hits with fixed charge-bin edges and refit. The central 68% interval uses the 16th and 84th percentiles of 1000 fits; Ref. [4] describes the general resampling principle.

Review N2-2: Use the frozen extended-profile metadata. Do not silently inherit newer code defaults. These full-statistics profiles are not train/test splits.

Reference version shared with Note 1

The run-4399 MCP functions were refitted on 22 October 2026 after correcting segment detection across missing trigger timestamps and excluding invalid clock groups. Note 1 documents the timing validation and remaining limitations.

For each topology, file and segment, the scaled MAD of the channel-192 versus scope-trigger line residuals must not exceed 1 ns before LOWESS [3]. The selection uses clock information, not SiPM timing residuals. Suspect event associations are excluded; no MCP rows are shifted.

All AWC-and-offset fits in this note use these refreshed references. The run-4076 slope fits, bootstrap replicas and measured-support matching definition are unchanged. The charge spectra use the updated reference sample.

AWC intervals still hold the reference fixed. They do not account for reference-fit uncertainty, clock-selection uncertainty or cross-run transfer.

Numerical version: `Q_Walk/results/2026-10-22_clock_quality`.

Sensitivity to the refreshed reference

The run-4076 input is fixed; these changes arise from the reference rerun, including its automatically determined charge selection and measured support.

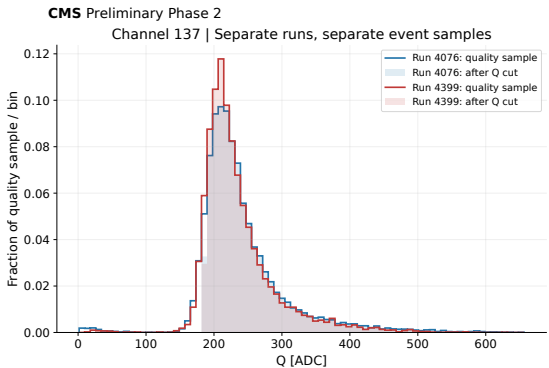
Channel	Previous AWC [μA]	Refreshed AWC [μA]
137	18.513	18.332
136	27.899	-4.240
151	7.649	9.354

Channel 136 has only 210 selected reference hits. Its sign change, and the 22% shift for channel 151, expose sensitivity not covered by slope-only errors. They are retained in the exploratory summary, not accepted as robust calibrations.

Sparse clock groups retain the existing line-only fallback when fewer than ten alignment inliers are available. A small residual scale in such a group does not independently validate its event association.

Review N2-6: Require reference stability and adequate clock statistics before defining the final channel-performance sample.

Charge spectra define the available comparison sample

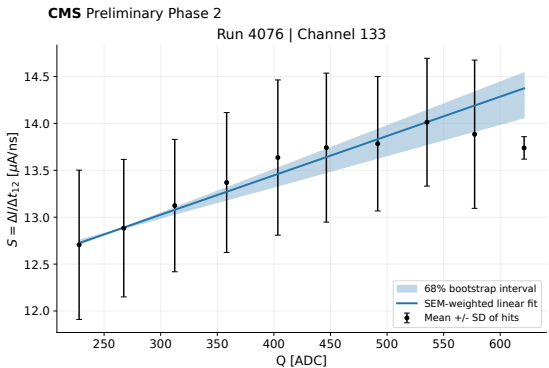


Channel 137 charge spectra are shown for the two independent runs. Lines show quality-selected hits; shaded spectra retain the additional charge-peak selection. Each run is normalized to its own quality-sample count.

The charge proxy is in ADC units, not MeV. Different populations and acceptances motivate restricting the curve alignment to jointly supported charge values.

Recovered histogram counts, not a new peak fit. Run-4076 profile selection starts at about 184.9 ADC for this channel.

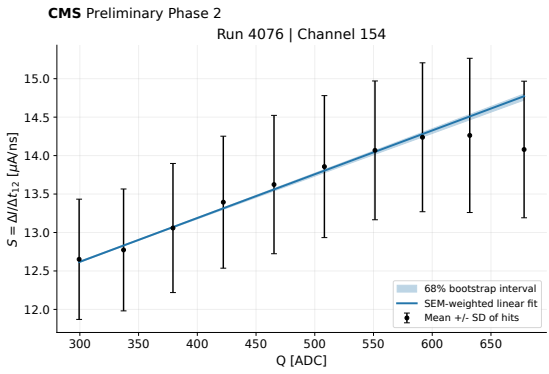
Channel 133: fitting the slope profile



Run 4076 channel-133 hits are grouped into ten charge bins. Points show mean S; bars show the SD of individual selected hits. The line is the saved SEM-weighted fit $c_0 + c_1 Q$.
The shaded interval is the central 68% bootstrap interval of the fitted mean curve. It is shown at its true size, without visual enlargement.

Extended-profile snapshot. SD describes hit spread; SEM and the fit band describe different statistical quantities.

Channel 154: the opposite end of bar 7

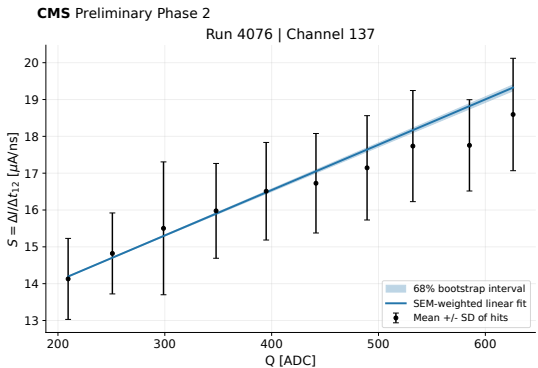


The same selection and profile procedure is applied to channel 154. The mean slope increases with charge, but the response is calibrated independently at each bar end. Bars show hit SD. The line uses SEM weights and the band propagates the saved bootstrap coefficient replicas, including their correlation.

Review N2-3: Assess linear-model adequacy separately from the statistical band; bootstrap spread is not a model-bias uncertainty.

Run 4076, extended-profile snapshot. No model refit is performed for the slide rendering.

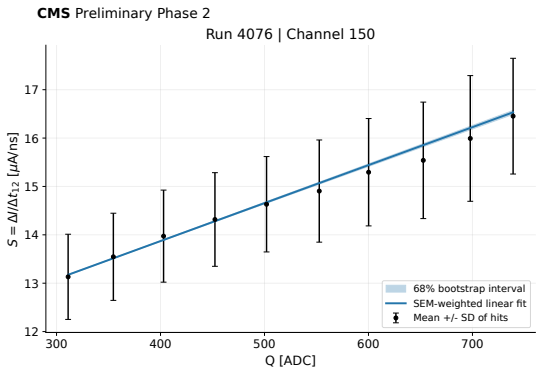
Channel 137: pulse slope and charge



The selected run-4076 sample contains 19 126 channel-137 hits. Ten profile points constrain a linear pulse-slope model over the observed charge range. The SD bars show event-to-event spread, while SEM weights determine the fit. The narrow shaded band is the conditional 68% bootstrap interval of the mean prediction, not a single-hit uncertainty.

Selected count and fit are read from the saved extended-profile metadata. No uncertainty band is magnified.

Channel 150: the opposite end of bar 8



Channel 150 is treated with the same profiling procedure. The charge-dependent slope provides the denominator of the proposed walk function; a separate amplitude A determines the correction strength. The line and shaded interval describe the fitted mean. Points and SD bars retain information on the spread within the selected charge bins.

Run 4076, extended-profile snapshot. Constant AWC does not mean a constant timing correction.

Determine AWC using only jointly measured support

The run-4399 reference is a fitted quadratic in inverse charge,

$$W_{\mathrm{MCP}}(Q) = a + b/Q + c/Q^2.$$

Let Ω be the intersection of the two outer profile-mean ranges, restricted to charge bins populated in both runs. Fit

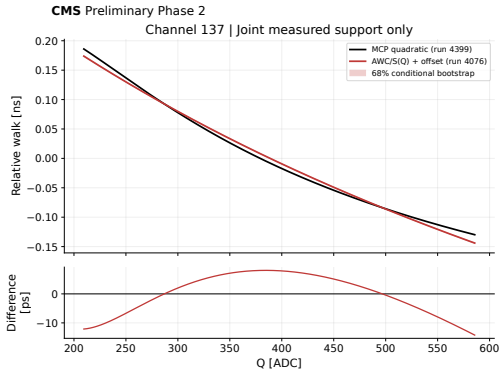
$$(A, B) = \arg\min_{A,B} \sum_{Q_k \in \Omega} \left[\frac{A}{c_0 + c_1 Q_k} + B - W_{\mathrm{MCP}}(Q_k) \right]^2 .$$

The saved implementation samples a uniform 3000-point charge grid before removing unsupported intervals. The offset B absorbs the arbitrary time zero.

This is uniform-charge shape matching, not an event-weighted likelihood. It uses neither extrapolated tails nor a single central anchor point.

Review N2-4: Compare weighting alternatives and propagate reference-fit uncertainty. The current objective is a documented choice, not a uniquely optimal estimator.

Channel 137: inverse-slope and MCP walk shapes



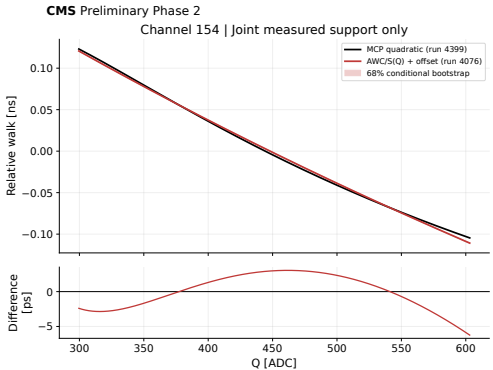
The black curve is the run-4399 MCP quadratic. The red curve is the run-4076 inverse-slope model after fitting AWC and a constant offset. Only jointly measured charge support is shown.

The lower panel is the difference between the two curves. Its RMS is 6.80 ps; this is not a single-event timing resolution.

The band varies the run-4076 slope fit while holding the MCP curve fixed.

AWC = 18.332 μA; conditional 16–84% interval [18.120, 18.607] μA. Both curves share one display offset.

Channel 154: a second shape-comparison example



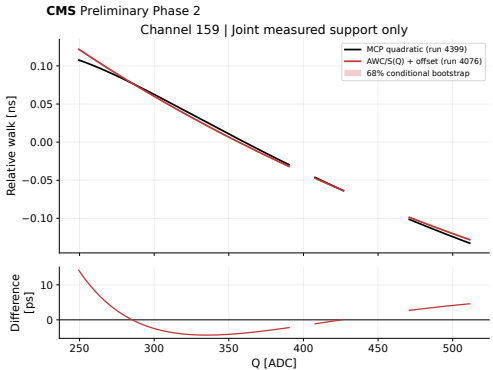
The same overlap-only objective determines AWC and the offset for channel 154. Both plotted curves are calibrated models, not matched events from different runs.

The lower panel directly displays the remaining shape mismatch, with no additional line subtracted.

Review N2-5: This replaces the ambiguous old residual illustration. It must be described as model closure, not a corrected-event resolution.

Conditional bootstrap band: 4076 slope replicas, refitted AWC and offset, fixed 4399 quadratic.

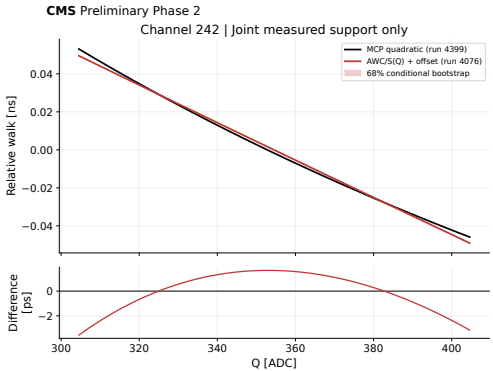
Channel 159: missing support is not extrapolation



Only intervals with populated profile bins in both runs contribute to the alignment. Breaks in the displayed curves mark excluded charge intervals. They do not represent discontinuities in either analytic model.
The lower panel uses exactly the same supported points. All displayed intervals participate in the common AWC-and-offset fit.

The stored comparison excludes 703 of the 3000 initial grid points for this channel. No line is drawn through those gaps.

Downstream example: channel 242

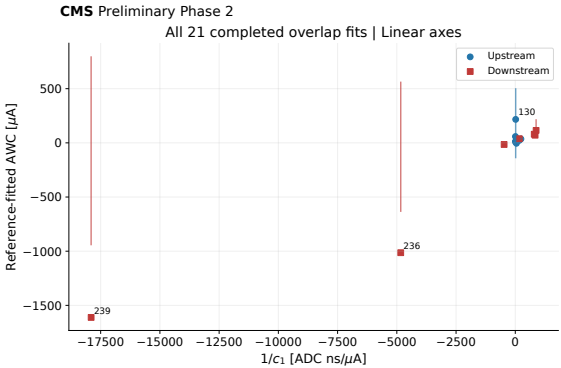


The method is also applied to a downstream channel. In run 4399, upstream bar 8 supplies the coincidence pair alongside the downstream target pair and MCP indicator.
In run 4076, the target downstream pair is required, without fixing an upstream pair. The curves use only their common measured support.

Review N2-1: Verify the transfer conditions and selection-induced changes for downstream channels as well as upstream channels.

The comparison is a per-channel shape calibration. The other module is not used as an internal timing reference to determine AWC here.

AWC versus inverse slope coefficient: all completed fits

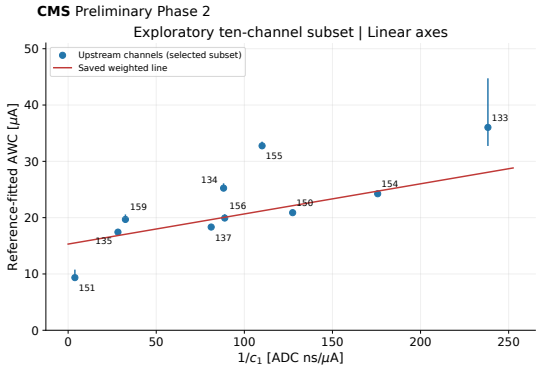


Each point is one channel, not one charge bin. Blue circles are upstream and red squares downstream. The x coordinate is $1/c_1$, the inverse coefficient of the linear $S(Q)$ fit; the y coordinate is reference-fitted AWC. Bars span the 16th–84th bootstrap percentiles. There are 21 completed fits; channel 238 has no measured overlap.

Review N2-6: Publish eligibility and failure criteria, including narrow support and poorly constrained channels.

Linear axes. Large intervals and negative coefficients are retained; this is not a final quality-selected performance sample.

A selected subset does not establish a universal law



The ten-channel upstream subset previously selected for discussion is shown on linear axes. Bars are conditional AWC bootstrap intervals; channel numbers identify the points. The saved weighted straight line has $\chi^2/\mathrm{ndf} = 538.6/8$ under its limited error model. It does not adequately describe the variation.

Review N2-6: The subset was selected interactively. It is exploratory, not an unbiased population-level validation.

The fit treats x as fixed and uses half the 16–84% AWC span as its symmetric uncertainty. Correlations and reference errors are omitted.

What the present uncertainties do and do not include

Displayed quantity	Meaning
Profile bars	SD of selected hits in each charge bin; not the uncertainty on the mean.
Fit weights	SEM of each profile mean; not the SD.
Slope band	16–84% range from 1000 bootstrap linear fits with fixed bin edges.
AWC bars / walk band	Propagate the slope replicas through a new AWC-plus-offset alignment to the same fixed reference curve.

Not yet included: run-4399 quadratic uncertainty, cross-run response transfer, selection/model choices and all shared-feature/target correlations.

Review N2-4: These are conditional statistical intervals, not total calibration uncertainties. A small curve mismatch is not a resolution estimate.

Result and next validation step

- A dual-threshold measurement provides a charge-dependent slope proxy; an inverse-slope function can be aligned to independent MCP walk curves.
- Restricting alignment to jointly measured charge support avoids determining AWC from extrapolation.
- Reference-calibrated AWC varies between channels. A straight line against $1/c_1$ is insufficient for the explored subset.
- The present calibration still uses the MCP to determine AWC. It does not establish reference-free AWC extraction by itself.
- An internal predictor must be tested on held-out bars/configurations with transfer and target uncertainties included before use without MCP.

The following backup explores such a predictor; it is not a final performance claim or a demonstration with cosmic rays.

Backup: an internal predictor for channel AWC

For $W = A/S$ and $S = c_0 + c_1 Q$,

$$-\frac{dW}{d\ln Q} = \frac{A\, c_1 Q}{S(Q)^2}.$$

If this local walk derivative is approximately shared between channels at a prescribed typical charge Q_*, a candidate model is

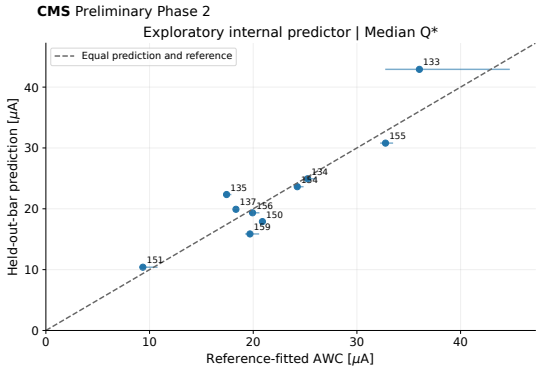
$$A_{\mathrm{pred}} = K\, \frac{S(Q_*)^2}{c_1 Q_*}.$$

Here Q_* is the median charge of the run-4076 quality- and peak-right-selected sample. It is computed once per channel; A_{pred} remains constant when correcting individual events.

The common K is learned from MCP-calibrated training channels. Predicting a new channel uses internal features, but the model development is reference-based.

Review N2-7: Shared local derivative is a hypothesis. Test operating-condition dependence and avoid selecting Q_* only after inspecting prediction scores.

Backup: predict a bar excluded from the model fit



Each point compares reference-fitted AWC with a prediction made after excluding that channel's entire bar from the fit of K. Ten upstream channels span six held-out bar groups.
The RMS prediction error is 3.22 μA; the RMS fractional error is 14.1%.
Horizontal intervals are conditional reference-AWC intervals, not predictive intervals.

Review N2-7: Small selected sample; feature/target correlations and full prediction uncertainties remain unpropagated.

Exploratory median-Q_* candidate. The dashed line indicates equality, not a fitted result. No blanket 10% precision claim is made.

References

[1] CMS Collaboration, *A MIP Timing Detector for the CMS Phase-2 Upgrade*, CERN-LHCC-2019-003, CMS-TDR-020 (2019). CERN document record.

[2] E. Albuquerque et al., *TOFHIR2: The readout ASIC of the CMS Barrel MIP Timing Detector* (2024). arXiv:2404.01208.

[3] W. S. Cleveland, *Robust Locally Weighted Regression and Smoothing Scatterplots*, J. Am. Stat. Assoc. **74** (1979) 829–836. doi:10.1080/01621459.1979.10481038.

[4] B. Efron, *Bootstrap Methods: Another Look at the Jackknife*, Ann. Statist. **7** (1979) 1–26. doi:10.1214/aos/1176344552.

These sources provide detector/readout context and methodological background. The run-specific selections, implementations and numerical results are those of this study, not measurements or prescriptions taken from these references.

Review queue before publication

ID	Decision or missing evidence
N2-1	Freeze cross-run operating conditions and justify calibration transfer.
N2-2	Document safe-t_2, threshold conversion and exact selected populations.
N2-3	Check linear-profile adequacy and selection/model sensitivity.
N2-4	Validate alignment objective and complete the uncertainty propagation.
N2-5	Keep model closure distinct from event timing-resolution validation.
N2-6	Freeze channel eligibility independently of observed trend; publish exclusions.
N2-7	Decide whether internal prediction remains backup or gains enough validation for a main result.

Sources: supplied `Note2.pdf`; extended profile snapshots and the dated measured-overlap study. All plotted tables and checksums accompany this draft.

Primary-source references are included. Official template, DP metadata and cover handling still require convener confirmation. This review slide is removed from a public release.